

SQL – Part 1

Running example: assumptions

- We use the following relation schemas:

Student(Sid, Sname, Major, Byear)

Course(Cno, Cname, Dept)

Enroll(Sid, Cno, Grade)

- Sid and Cno are foreign keys in Enroll referencing Sid and Cno in Student and Course, respectively.
- We will also assume that Student has m tuples, Enroll has n tuples, and Course has k tuples.

SQL queries – basic form

- The simplest form of a SQL query is as follows:

SELECT list of components of tuple variables
FROM list of tuple variables associated with relations
WHERE condition on components of tuple variables;

Example of a simple SQL query

- Consider the query “Find the sids and names of students who major in CS.”
- In SQL, the formulation is,

```
SELECT S.Sid, S.Sname  
FROM   Student S  
WHERE  S.Major = 'CS';
```

Semantics of a basic SQL query

- In the SQL query

```
SELECT S.Sid, S.Sname  
FROM Student S  
WHERE S.Major = 'CS';
```

S is a tuple variable that ranges over all tuples in the *Student* relation. So, if this relation has 10 tuples, then *S* ranges over these 10 tuples.

- The semantics of this query is

```
for each tuple S in Student  
  if S.Major = 'CS'  
    then output (S.Sid, S.Sname)
```


Semantics (continued)

- It is important to realize that, when a tuple variable is associated with a relation, then that tuple variable is assigned, one-at-a-time, to each tuple in that relation.

- E.g., if **Student** is

<u>Sid</u>	Sname	Major	Byear	
s1	John	CS	1990	← S
s2	Ellen	Math	1995	← S
s3	Eric	CS	1990	← S
s4	Ann	Biology	2001	← S

then **S** is assigned, one at the time, to these four tuples.

- The order in which **S** is assigned to these tuples is not pre-determined since the **Student** relation is a set!

Semantics (continued)

- For each tuple assigned to the variable **S**, the (boolean) condition in the **WHERE** clause is checked.
- If the condition is true for that tuple, then the attributes of that tuple specified in the **SELECT** clause will be output.
- If the condition is false for that tuple, no attributes of that tuple are output.

SQL has bag semantics

- Consider the query “Find the majors of students.”
- In SQL:

```
SELECT S.Major  
FROM Student S
```
- Semantics: **for each tuple S in Student**
output (S.Major)
- Output on our sample database:
- Observe that ‘CS’ appears twice: the result of this query is a **bag** (**multiset**)

Major
CS
Math
CS
Biology

Enforcing set semantics

- To coerce the result of the query into a set, use the **DISTINCT** clause:

```
SELECT DISTINCT S.Major  
FROM Student S
```

- The result is the set

Major
CS
Math
Biology

Time-complexity of DISTINCT

- The time-complexity of the coercion from the bag into the set is at least linear in the size of the bag.
- If **hashing** is used, it can be done during the output generation.
- One can also use **sorting** as a technique to accomplish the coercion. In this case, the time-complexity is $O(n \log n)$, where n is the number of tuples in the bag.

EXPLAIN in PostgreSQL

- In PostgreSQL, it is possible to see how the system implements the coercion.
- To see this, issue the **EXPLAIN** command:

```
EXPLAIN SELECT DISTINCT S.Major  
FROM Student S;
```

- The system will generate an **access plan** that indicates whether hashing or sorting is used.
- In general, **EXPLAIN** is very useful to understand how an SQL query is processed by the system.

EXPLAIN ANALYZE in PostgreSQL

- In PostgreSQL, it is also possible to see prediction and actual time and space complexity associated with a query.
- To see this, issue the **EXPLAIN ANALYZE** command:

```
EXPLAIN ANALYZE SELECT DISTINCT S.Major  
FROM Student S;
```


Renaming attributes

- In SQL, it is possible to give names to attributes that appear in a SELECT clause using the **AS** clause:

```
SELECT S.Sid AS Identifier, S.Sname AS Name  
FROM Student S  
WHERE S.Major = 'CS';
```

- This query returns the relation

Identifier	Name
s1	John
s3	Eric

Ordering the output

- In SQL, it is possible to sort the output on the basis of order relation that exists on attribute domains
- This can be done with the **ORDER BY** clause

```
SELECT S.sid, S.sname  
FROM Student S  
WHERE S.Major = 'CS'  
ORDER BY sname;
```

sid	sname
s3	Eric
s1	John

Ordering the output (by column)

- The following queries are equivalent

```
SELECT S.sid, S.sname  
FROM Student S  
WHERE S.Major = 'CS'  
ORDER BY sname;
```

```
SELECT S.sid, S.sname  
FROM Student S  
WHERE S.Major = 'CS'  
ORDER BY 2;
```

Randomly ordering the output

- In SQL, it is possible to randomly order the output of a query
- This can be done with the `ORDER BY RANDOM()` clause
- The following executions of the same queries may not give the same result

```
SELECT *  
FROM Student S  
ORDER BY RANDOM();
```

```
SELECT *  
FROM Student S  
ORDER BY RANDOM();
```


Queries involving multiple relations

- Consider the query “Find the names of students along with the cnos of the courses in which they received a B grade”:

```
SELECT S.Sname, E.Cno  
FROM   Student S, Enroll E  
WHERE  S.Sid = E.Sid AND E.Grade = 'B';
```

- We have **two** tuple variables:
 - S** ranges over the tuples in **Student**;
 - E** ranges over the tuples in **Enroll**.
- In the **WHERE** clause, conditions are checked on attribute values of tuples assigned to these variables.

Queries on multiple relations (c'ed)

- SQL query:

```
SELECT S.Sname, E.Cno  
FROM   Student S, Enroll E  
WHERE  S.Sid = E.Sid AND E.Grade = 'B';
```

- Semantics:

```
for each tuple S in Student  
  for each tuple E in Enroll  
    if S.Sid = E.Sid and E.Grade = 'B'  
      then output (S.Sname, E.Cno)
```

- Notice the double-nested for-loop!

Queries on multiple relations (c'ed)

- SQL query:

```
SELECT S.Sname, C.Cno  
FROM   Student S, Enroll E  
WHERE  S.Sid = E.Sid AND E.Grade = 'B'
```

- The condition $S.Sid = E.Sid$ is called a **join condition**: it verifies a condition between two variables
- The condition $E.Grade = 'B'$ is called a **constant comparison** condition: it verifies a condition on a single variable.

Time-complexity considerations

- If Student has m tuples and Enroll has n tuples, then the if statement

if S.Sid = E.Sid and E.Grade = 'B'

is evaluated $m \times n$ times.

- This essentially quadratic time complexity compromises the scalability of this simple query for large values of m and n .

Sub-queries in the FROM clause

- The previous query can also be formulated as follows:

```
SELECT S.Sname, C.Cno
FROM Student S, (SELECT E.Sid, E.Cno
                  FROM Enroll E
                  WHERE E.Grade = 'B') C
WHERE S.Sid = C.Sid;
```

- There are 2 **global** variables (**S**, **C**) and 1 **local** variable (**E**):
 - **S** ranges over the tuples in **Student**;
 - **C** ranges over the tuples in the relation determined by the **sub-query** (the part between the brackets);
 - **E** ranges over the tuples in **Enroll**.

Subqueries (continued)

- We conclude the following:
 - In SQL, it is permissible to have tuple variables that range over relations, but also tuple variables that range over the tuples of a relation computed by a subquery.
 - This works, because an SQL query always returns a relation.
 - Subqueries are allowed to occur in the **FROM** clause of a query or another subquery.

Performance improvement

- Query structure:

```
SELECT S.Sname, C.Cno  
FROM Student S, (subquery) C  
WHERE S.Sid = C.Sid
```
- Semantics:

```
for each tuple S in Student  
  for each tuple C in output (subquery)  
    if S.Sid = C.Sid  
      then output (S.Sname, C.Cno)
```
- If the subquery yields s tuples, then the if-then statement will be executed $m \times s$ times instead of $m \times n$.
- This can be a considerable gain if s is small compared to m .

- Of course we must also add the time to compute the relation generated by the subquery. This can be done using a linear scan requiring n operations.
- So the total time is $n + m \times s$ instead of $m \times n$.

Multiple tuple variables ranging over the same relation

- Consider the query “Find the name of each student enrolled in at least **two** courses.”
- Since a tuple variable ranging over a relation is only assigned to a single tuple of that relation, it can never hold information about two or more tuples.
- However, if we have two or more tuple variables ranging over the same relation, then this becomes possible.

Multiple tuple variables ranging over the same relation: example

- We want to express the query “Find the sid and name of each student enrolled in at least **two** courses.”
- In SQL, this query can be formulated using two variables ranging over the **Enroll** relation:

```
SELECT DISTINCT S.Sid, S.Sname  
FROM   Student S, Enroll E1, Enroll E2  
WHERE  S.Sid = E1.Sid AND S.Sid = E2.Sid AND E1.Cno <> E2.Cno;
```

- **<>** is the symbol for not equal, i.e., \neq .
- Why do we need DISTINCT?

Multiple tuple variables ranging over the same relation: example (continued)

- Semantics:

```
for each tuple S in Student
  for each tuple E1 in Enroll
    for each tuple E2 in Enroll
      if S.Sid = E1.Sid and S.Sid = E2.Sid and E1.Cno ≠ E2.Cno
        then output (S.Sid, S.Sname)
```

- Performance: $m \times n^2$.
- Later, we shall see that there are far more efficient ways to write and solve this query.

Queries using set operations

- We now consider queries involving the set operations **union**, **intersection**, and **difference**:

Operation	Math notation	SQL notation
Union	\cup	UNION
Intersection	\cap	INTERSECT
Difference	$-$	EXCEPT

Set operations: examples

R

A	B
1	1
1	2
3	2
3	1

S

A	B
1	3
3	1
3	2
2	2
2	3

$R \cup S$

A	B
1	1
1	2
3	2
3	1
1	3
2	2
2	3

$R \cap S$

A	B
3	2
3	1

$R - S$

A	B
1	1
1	2

$S - R$

A	B
1	3
2	2
2	3

Notice that $R \cup S = S \cup R$ and $R \cap S = S \cap R$.
However, in general, $R - S \neq S - R$.

Union

- Consider the query “Find the sids and names of all students who major in CS **or** who major in Math.”
- In SQL,

```
(SELECT S.Sid, S.Sname  
FROM Student S  
WHERE S.Major = 'CS')  
UNION  
(SELECT S.Sid, S.Sname  
FROM Student S  
WHERE S.Major = 'Math');
```


Union (alternative)

- Consider again “Find the sids and names of all students who major in CS **or** who major in Math.”
- In SQL,

```
SELECT S.Sid, S.Sname  
FROM   Student S  
WHERE  S.Major = 'CS' OR S.Major = 'Math';
```

Intersection

- Consider the query “Find the sids of all students who are enrolled in course c1 **and** course c2.”
- In SQL,

```
(SELECT E.Sid  
FROM   Enroll E  
WHERE  E.Cno = 'c1')  
INTERSECT  
(SELECT E.Sid  
FROM   Enroll E  
WHERE  E.Cno = 'c2');
```


Wrong alternative for intersection

- Consider again “Find the sids of all students who major in c_1 **and** c_2 .”
- What about the following SQL query?

```
SELECT E.Sid  
FROM   Enroll E  
WHERE  E.Cno = 'c1' AND E.Cno = 'c2';
```

- The result of this query is the **empty set** since, for no tuple from **Enroll** assigned to the variable **E**, **E.Cno** can be equal to both **c_1** and **c_2** at the same time!

Difference

- Consider the query “Find the sids of all students who are enrolled in **no** courses.”
- In SQL,

```
(SELECT S.Sid
FROM Student S)
EXCEPT
(SELECT E.Sid
FROM Enroll E);
```
- The first subquery yields the set of all student sids.
- The second subquery yields the set of the sids of all students enrolled in some course.
- The difference between both sets yields the result.

Performance considerations

- UNION, INTERSECTION, and DIFFERENCE can be implemented using hashing or using sorting.
- Therefore, these operations can be implemented in linear time $O(m + n)$ or in $O(n \log n + m \log m)$, where m and n are the respective sizes of the relations involved in these operations.

Semantics of set operations in SQL

- In SQL, the result of the **UNION**, **INTERSECT**, and **EXCEPT** operations are **sets**, even if the inputs to these operations are bags.
- So, the result of

```
(SELECT E.Sid  
FROM Enroll E)  
UNION  
(SELECT E.Sid  
FROM Enroll E)
```

is the same as the result of

```
SELECT DISTINCT E.Sid  
FROM Enroll E
```
- To retain **bag** semantics, use **UNION ALL**, **INTERSECT ALL**, and **EXCEPT ALL**.