

Aggregate Functions and Data Partitioning

Collections and aggregate functions

- A **collection** is a grouping of some variable number of data items (possibly zero)
- Usually the data items in a collection are of the **same type**
- **Aggregate functions** are functions that apply to collections, i.e., they consider **all** these data items in these collections
- Applied to a collection, an aggregate function returns a single value



Examples of collections

- sets, multisets, dictionaries (maps), relations
- vectors, lists, arrays, series
- data structures: stacks, queues, hash tables, trees, graphs

Aggregate functions on unordered collections

- We will restrict ourselves to aggregate functions on sets, bags, and relations
- COUNT (we will often use the notation $|A|$ instead of COUNT(A))
- SUM, AVERAGE, MIN, and MAX
- $\text{COUNT}(\{a,b,c\}) = 3$; $\text{COUNT}(\{\{a,c,c,b,b,b\}\}) = 6$
- $\text{SUM}(\{1, 4, 7\}) = 12$; $\text{SUM}(\{\{1,1,1,1,1\}\}) = 5$
- $\text{AVG}(\{1,4,7\}) = 4$; $\text{AVG}(\{\{1,1,1,1,1\}\}) = 1$
- $\text{MIN}(\{1,4,7\}) = \text{MIN}(\{\{1,1,1,1,1\}\}) = 1$
- $\text{MAX}(\{1,4,7\}) = 7$; $\text{MAX}(\{\{1,1,1,1,1\}\}) = 1$
- $\text{MIN}(\{\text{John}, \text{Eric}, \text{Ann}\}) = \text{Ann}$
- These aggregate functions are supported in SQL



Applications of aggregate functions

- Data analytics
- Formulating complex queries
- Facilitating efficient query evaluation

Counting the size of a set in SQL

- Let $R(A_1, \dots, A_n)$ be a relation.
- Then $|R|$ can be obtained as follows:

```
SELECT COUNT(*)  
FROM   R r;
```

- Alternatively,

```
SELECT COUNT(1)  
FROM   R r;
```

- Of course we can restrict the **COUNT** function to apply to a subset of R by applying a **WHERE** clause.

COUNT examples

- “Find the number of courses in which the student with sid 10 is enrolled.”

```
SELECT COUNT(*)  
FROM   Enroll E  
WHERE  E.sid = 's10';
```

- “Find the number of students who are not enrolled in any CS course.”

```
SELECT COUNT(*)  
FROM   Student S  
WHERE  S.Sid NOT IN (SELECT E.Sid  
                     FROM   Enroll E, Course C  
                     WHERE  E.Cno = C.Cno AND C.Dept = 'CS');
```

COUNT Example

- Let R and S be two relations, then the following query will return $|R \times S| = |R||S|$, i.e. the size of the cartesian (cross) product of R and S .

```
SELECT COUNT(*)  
FROM R r, S s;
```

- $R \times S = \{(r, s) \mid r \in R \wedge s \in S\}$

COUNT DISTINCT

R

A	B
a	1
a	2
b	1

SELECT COUNT(r1.A) AS Total
FROM R r1, R r2



Total
9

SELECT COUNT(DISTINCT r1.A) AS Total
FROM R r1, R r2



Total
2

Simulating COUNT with SUM

- The following SQL query uses the SUM aggregate function to determine the size of R, provided $R \neq \emptyset$

```
SELECT SUM(1)  
FROM   R r;
```

- The bag that is generated by the query is $\{\{1, \dots, 1\}\}$ containing as many 1's as there are tuples r in R

Example

A
a
b
c

- Consider the relation R
- $|R| = 3$
- COUNT applies to $\{a, b, c\}$
- SUM applies to $\{1, 1, 1\}$
- If you write the SQL query

```
SELECT SUM(2)
FROM R r;
```

the result will be $SUM(\{2, 2, 2\}) = 6$.

Caveat: empty collection

- Consider the relation R
- $|R| = 0$
- COUNT applied to {} gives 0
- SUM applied to {} gives NULL

A

```
SELECT COUNT(1)  
FROM R r;
```



0

```
SELECT SUM(1)  
FROM R r;
```



NULL

MIN and MAX aggregate functions

- **MIN** returns the smallest data item in the bag to which it applies.
- **MAX** returns the largest data item in the bag to which it applies.
- Data items can come from any ordered basic domain: integer, float, text
- A more general **MIN** function can be simulated using **<= ALL**

```
SELECT DISTINCT r.A1,...,r.An  
FROM   R r  
WHERE (r.A1,...,r.An) <= ALL (SELECT r1.A1,...,r1.An  
                             FROM   R r1);
```

- **MAX** can be simulated using **>= ALL**
- However, there is a problem if MIN (MAX) is applied to an empty set.

CAVEAT: aggregate functions on empty set

- Except for COUNT, SQL aggregate functions return a NULL value when applied to an empty set (or bag).
- Assume R is the empty relation

A

- Then `SELECT MIN(r.A) AS smallest FROM R r`
returns the relation

smallest
NULL
- However `SELECT r.A AS smallest FROM R r WHERE r.A <= ALL(SELECT r1.A FROM R r1)`
returns the empty relation

smallest

Partitioning and counting

- “Determine for each student the number of courses taken by that student.”

Enroll


Sid	Cno
S1	C1
S2	C1
S1	C2
S3	C3
S3	C4
S1	C3



Sid	No_Courses
S1	3
S2	1
S3	2

Partition and map count function

- (1) First, **Partition** the Enroll table into cells (blocks) wherein each cell contains all the tuples that have a common sid value.
- (2) Next, **Map** the **COUNT** function over these cells.



Partition and map COUNT in SQL

- (1) The **GROUP BY** map **COUNT** method
- (2) The user-defined **COUNT FUNCTION** method
- (3) The **SELECT COUNT-expression** method

The GROUP BY map COUNT method

Map COUNT phase → `SELECT E.Sid, COUNT(*) AS No_Courses
FROM Enroll E`
Partition phase → `GROUP BY(E.Sid)`

Partition phase: the GROUP BY operator places each tuple E into the cell identified by its E.Sid value

Map COUNT phase: the COUNT function is mapped over the cells identified by the different possible E.sid values

Group by visualization

```
SELECT State, COUNT(*)  
FROM facebook  
GROUP BY State
```

facebook

Name	State	# of Friends
Matt	CA	300
Lisa	CA	500
Jeff	CA	600
Sarah	FL	400

final table

State	COUNT(*)
CA	0
FL	0

<https://dataschool.com/how-to-teach-people-sql/how-sql-aggregations-work/>

The user-defined COUNT FUNCTION method

- “Define a function with input parameter a student sid and as output the number of courses taken by that student.

```
CREATE FUNCTION NumberOfCourses (s TEXT) RETURNS bigint  
AS $$
```

```
    SELECT COUNT(*)  
    FROM   Enroll E  
    WHERE E.Sid = s;  
$$ LANGUAGE SQL;
```

 Map COUNT phase

- Then execute the SQL query

```
SELECT S.Sid, NumberOfCourses(S.Sid) AS No_Courses  
FROM   Student S;
```

 Partition phase

Student

Sid
S1
S2
S3
S4

Enroll

Sid	Cno
S1	C1
S1	C2
S1	C3
S2	C3
S3	C3
S3	C4

SELECT S.Sid, NumberOfCourses(S.Sid) AS No_Courses
FROM Student S;



Sid	No_Courses
S1	3
S2	1
S3	2
S4	0

student s4 takes no courses



The SELECT COUNT-expression method

```
SELECT S.sid, (SELECT COUNT(E.Cno) AS NumberCourses
               FROM   Enroll E
               WHERE  E.Sid = S.Sid)
FROM   Student S
```

← Map phase

← Partition phase

- Observe that the subquery identified by S.Sid appears in the outer SELECT clause.
- The COUNT of the result of this subquery is then delivered as a value in the outer SELECT clause.
- Notice that this expression must appear between parentheses.
- The output of this query is the same as that on the previous slide.

Example query

- “Find the sid of each student who take the most courses.”
- Using the GROUP BY method:

WITH

```
NumberOfCoursesbyStudent AS (SELECT E.Sid, COUNT(E.Cno) As NumberOfCourses
                             FROM   Enroll E
                             GROUP BY(E.Sid))
```

[illegible]

Example query

- “Find the sid of each student who takes the most courses.”
- Using the COUNT expression method:

```
SELECT S.Sid
FROM Student S
WHERE (SELECT COUNT(E.cno)
      FROM Enroll E
      WHERE E.sid = S.sid) >= ALL (SELECT (SELECT COUNT(E.cno)
                                           FROM Enroll E
                                           WHERE E.sid = S1.sid)
                                   FROM Student S1);
```

The COUNT-bug of GROUP BY

- The result of the following 2 queries is the same. Notice that there is a bug since, if a student sid takes no courses, then (sid, 0) does not appear in the output.

```
SELECT E.Sid, COUNT(E.Cno)
FROM   Enroll E
GROUP BY (E.Sid)
```

```
SELECT S.Sid, Count(E.Cno)
FROM   Student S, Enroll E
WHERE S.Sid = E.Sid
GROUP BY(S.Sid)
```

- These two queries give the same result and exhibit the COUNT bug: the tuple (s4,0) does not appear in the result.

Fixing the COUNT-bug

- To fix the COUNT-bug we need to add the (s,0) pair if student with sid s takes no courses. This can be done with the **UNION** operator.

```
(SELECT E.Sid, COUNT(E.Cno) AS No_Courses
FROM   Enroll E
GROUP BY (E.Sid))
UNION
(SELECT S.Sid, 0 AS No_Courses
FROM   Student S
WHERE  S.Sid NOT IN (SELECT E.Sid
                     FROM   Enroll E))
```

Partitioning on different dimensions

Enroll

Sid	Cno
S1	C1
S1	C2
S1	C3
S2	C3
S3	C3
S3	C4

```
SELECT COUNT(*)  
FROM Enroll  
GROUP BY ()
```



6

```
SELECT E.Sid, COUNT(*)  
FROM Enroll E  
GROUP BY (E.Sid)
```



E.sid	
S1	3
S2	1
S3	2

```
SELECT E.Sid, E.Cno, COUNT(*)  
FROM Enroll E  
GROUP BY (E.Sid, E.Cno)
```



E.Sid	E.Cno	
S1	C1	1
S1	C2	1
S1	C3	1
S2	C3	1
S3	C3	1
S3	C4	1

What can appear in the GROUP BY clause?

- Answer: any valid expression over the tuples in the FROM clause.

```
SELECT ...  
FROM R1 t1,...,Rn tn  
WHERE ...  
GROUP BY(expression(t1,...,tn))
```

- Partition: there will be as many cells in the partition as there are different values for **expression(t₁,...,t_n)**

Example: expressions in GROUP BY

S

X	Y
2	3
1	3
2	1
0	3

```
SELECT s.x + s.y AS sum, COUNT(*) as cell_size  
FROM   S s  
GROUP BY (s.x + s.y)
```



sum	cell_size
3	2
4	1
5	1

Notice that there are 4 tuples in R assigned to variable s but only 3 different s.x+s.y values: 3, 4, 5. Thus there are only 3 cells in the partition.

Example: expression in GROUP BY

Person

Pid	Age
p1	10
p2	9
p3	12
p3	9

```
SELECT  p.age > 10 AS OlderThanTen, COUNT(*)  
FROM    Person p  
GROUP BY (p.age > 10);
```



OlderThanTen	Count
f	3
t	1

Restrictions on SELECT clause in GROUP BY query

- In a GROUP BY query, the SELECT clause may only contain **aggregate expressions** that returns a **single value** for each cell of the partition induced by the GROUP BY clause.

```
SELECT AggregateExpression(cell(expression(t1,...,tn)), ...  
FROM   R t1,...,tn  
WHERE  condition(t1,...,tn)  
GROUP BY (expression(t1,...,tn))
```


Aggregate expressions in SELECT clause

S

X	Y
2	3
1	3
2	1
0	3

```
SELECT s.x+s.y AS sum,  
       SUM(s.x*s.y) AS  
       sum_of_products
```

```
FROM S s
```

```
GROUP BY (s.x+s.y)
```



sum	sum_of_products
3	2 = (2 * 1 + 0 * 3)
4	3 = (1*3)
5	6 = (2*3)

Aggregate expressions in SELECT clause

- The following query will raise an error since s.x is **not necessarily unique** in a cell defined by s.x+s.y values

S

X	Y
2	3
1	3
2	1
0	3

```
SELECT s.x  
FROM S s  
GROUP BY (s.x+s.y)
```

The HAVING clause in GROUP BY queries

- The HAVING clause in a GROUP BY selects those cells from the partition induced by the GROUP BY clause that satisfy an **Aggregate Condition**.
- Only those cells are passed onto the SELECT clause.

```
SELECT AggregateExpression(Cell(expression(t1,...,tn)), ...  
FROM   R1 t1, ..., R tn  
WHERE  condition(t1,...,tn)  
GROUP BY (expression(t1,...,tn))  
HAVING AggregateCondition(Cell(expression(t1,...,tn)))
```

Example: HAVING clause

- “For each student who majors in CS determine the number of courses taken by that student, provided that this number is at least 2.”

```
SELECT E.Sid, COUNT(E.Cno)
FROM   Enroll E, Student S
WHERE  E.Sid = S.Sid AND S.Major = 'CS'
GROUP BY (E.Sid)
HAVING COUNT(E.Cno) ≥ 2;
```

Simulating HAVING clause with user-defined functions in WHERE

- “For each student who majors in CS, determine the number of courses taken by that student, provided that this number is at least 3.”
- The HAVING condition can be simulated in the WHERE clause with user-defined functions.

```
SELECT S.Sid AS Sid, NumberOfCourses(S.Sid)
FROM   Student S
WHERE  S.major = 'CS' AND NumberOfCourses(S.Sid) ≥ 3
```

Spreadsheet (Data Cube)

Sid	Cno
S1	C1
S1	C2
S1	C3
S2	C1
S2	C2
S3	C2
S4	C1

	c1	c2	c3	sum(sid)	
s1	1	1	1	3	
s2	1	1	0	2	
s3	0	1	0	1	
s4	1	0	0	1	
sum(cno)	3	3	1		
				7	sum()

GROUPING sets

- It may be desirable to simultaneously generate different partitions and then apply an aggregation
- This is supported in SQL via GROUPING sets.

Sid	Cno
S1	C1
S1	C2
S2	C1
S2	C2
S3	C1



Sid	Cno	Count
S1	*	2
S2	*	2
S3	*	1
*	C1	3
*	C2	2

```
SELECT Sid, Cno, COUNT(*)  
FROM Enroll  
GROUP BY  
GROUPING SETS((Sid),(Cno))
```