Aggregate Functions and Data Partitioning

Collections and aggregate functions

- A collection is a grouping of some variable number of data items (possibly zero)
- Usually the data items in a collection are of the same type
- Aggregate functions are functions that apply to collections, i.e., they consider all these data items in these collections
- Applied to a collection, an aggregate function returns a single value

Examples of collections

- sets, multisets, dictionaries (maps), relations
- vectors, lists, arrays, series
- data structures: stacks, queues, hash tables, trees, graphs

Aggregate functions on unordered collections

- We will restrict ourselves to aggregate functions on sets, bags, and relations
- COUNT (we will often use the notation |A| instead of COUNT(A))
- SUM, AVERAGE, MIN, and MAX
- COUNT({a,b,c}) = 3; COUNT({{a,c,c,b,b,b}}) = 6
- SUM({1, 4, 7}) = 12; SUM({{1,1,1,1,1}}) = 5
- AVG({1,4,7}) = 4; AVG({{1,1,1,1,1}}) = 1
- MIN({1,4,7}) = MIN({{1,1,1,1,1}}) = 1
- MAX({1,4,7}) = 7; MAX({{1,1,1,1,1}}) = 1
- MIN({John,Eric,Ann}) = Ann
- These aggregate functions are supported in SQL

Applications of aggregate functions

- Data analytics
- Formulating complex queries
- Facilitating efficient query evaluation

Counting the size of a set in SQL

- Let R(A1,…,An) be a relation.

- Then |R| can be obtained as follows:

    SELECT COUNT(*)
    FROM   R r;

- Alternatively,

    SELECT COUNT(1)
    FROM   R r;

- Of course we can restrict the COUNT function to apply to a subset of R by applying a WHERE clause.

COUNT examples

- "Find the number of courses in which the student with sid 10 is enrolled."

> SELECT COUNT(*)
> FROM   Enroll E
> WHERE E.sid = 's10';

- "Find the number of students who are not enrolled in any CS course."

SELECT COUNT(*)
FROM   Student S
WHERE S.Sid NOT IN (SELECT E.Sid
                    FROM   Enroll E, Course C
                    WHERE E.Cno = C.Cno AND C.Dept = 'CS');

COUNT Example

- Let R and S be two relations, then the following query will return $|R \times S| = |R||S|$, i.e. the size of the cartesian (cross) product of R and S.

SELECT COUNT(*)
FROM R r, S s;

- $R \times S = \{(r, s) \mid r \in R \land s \in S\}$

COUNT DISTINCT

R

A	B
a	1
a	2
b	1

SELECT COUNT(r1.A) AS Total
FROM   R r1, R r2

Total
9

SELECT COUNT(DISTINCT r1.A) AS Total
FROM   R r1, R r2

Total
2

Simulating COUNT with SUM

- The following SQL query uses the SUM aggregate function to determine the size of R, provided R $\neq$ $\emptyset$

SELECT SUM(1)
FROM   R r;

- The bag that is generated by the query is $\{\{1,...,1\}\}$ containing as many 1's as there are tuples r in R

Example

A
a
b
c

- Consider the relation R
- $|R| = 3$

- COUNT applies to {{a,b,c}}
- SUM   applies to {{1, 1, 1}}
- If you write the SQL query

        SELECT SUM(2)
        FROM   R r;

  the result will be SUM({{2,2,2}}) = 6.

Caveat: empty collection

- Consider the relation R

A

- $|R| = 0$

- COUNT applied to {} gives 0
- SUM   applied to {} gives NULL

```
SELECT COUNT(1)
FROM    R r;
```

➡️

0

```
SELECT SUM(1)
FROM    R r;
```

➡️

NULL

MIN and MAX aggregate functions

- **MIN** returns the smallest data item in the bag to which it applies.
- **MAX** returns the largest data item in the bag to which it applies.
- Data items can come from any ordered basic domain: integer, float, text
- A more general MIN function can be simulated using <= ALL

```
SELECT DISTINCT r.A1,...,r.An
FROM   R r
WHERE (r.A1,...,r.An) <= ALL (SELECT r1.A1,...,r1.An
                             FROM   R r1);
```

- **MAX** can be simulated using >= ALL
- However, there is a problem if MIN (MAX) is applied to an empty set.

CAVEAT: aggregate functions on empty set

- Except for COUNT, SQL aggregate functions return a NULL value when applied to an empty set (or bag).
- Assume R is the empty relation

A

- Then   SELECT MIN(r.A) AS smallest FROM R r

  returns the relation

smallest
NULL

- However SELECT r.A AS smallest FROM R r WHERE r.A <= ALL(SELECT r1.A FROM R r1)

  returns the empty relation

smallest

Partitioning and counting

- "Determine for each student the number of courses taken by that student."

Enroll
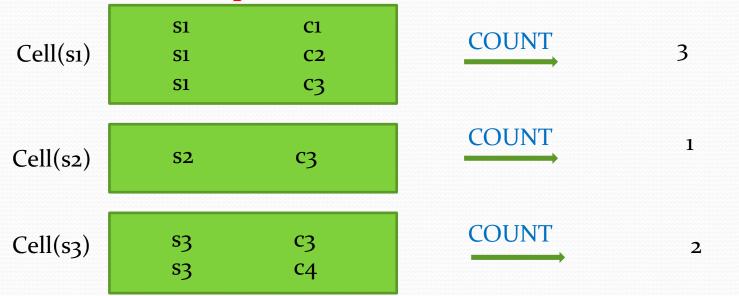
Sid	Cno
S1	C1
S2	C1
S1	C2
S3	C3
S3	C4
S1	C3

Sid	No_Courses
S1	3
S2	1
S3	2

Partition and map count function

- (1) First, Partition the Enroll table into cells (blocks) wherein each cell contains all the tuples that have a common sid value.

- (2) Next, Map the COUNT function over these cells.

Cell(s1)
s1	c1
s1	c2
s1	c3

COUNT → 3

Cell(s2)
s2	c3

COUNT → 1

Cell(s3)
s3	c3
s3	c4

COUNT → 2

Partition and map COUNT in SQL

- (1) The GROUP BY map COUNT method

- (2) The user-defined COUNT FUNCTION method

- (3) The SELECT COUNT-expression method

The GROUP BY map COUNT method

Map COUNT phase ➡ SELECT E.Sid, COUNT(*) AS No_Courses
FROM   Enroll E
Partition phase ⟶ GROUP BY(E.Sid)

Partition phase: the GROUP BY operator places each tuple E into the cell identified by its E.Sid value

Map COUNT phase: the COUNT function is mapped over the cells identified by the different possible E.sid values

Group by visualization

SELECT State, COUNT(*)
FROM facebook
GROUP BY State

facebook

Name	State	# of Friends
Matt	CA	300
Lisa	CA	500
Jeff	CA	600
Sarah	FL	400

final table

State	COUNT(*)
CA	0
FL	0

The user-defined COUNT FUNCTION method

- "Define a function with input parameter a student sid and as output the number of courses taken by that student.

```
CREATE FUNCTION NumberOfCourses (s TEXT) RETURNS bigint
AS $$
       SELECT COUNT(*)                                    ⟵ Map COUNT phase
       FROM    Enroll E
       WHERE E.Sid = s;
    $$ LANGUAGE SQL;
```

- Then execute the SQL query

```
SELECT  S.Sid, NumberOfCourses(S.Sid) AS No_Courses   ⟵ Partition phase
FROM    Student S;
```

Student

Sid
s1
s2
s3
s4

Enroll

Sid	Cno
s1	c1
s1	c2
s1	c3
s2	c3
s3	c3
s3	c4

SELECT S.Sid, NumberOfCourses(S.Sid) AS No_Courses
FROM   Student S;

Sid	No_Courses
s1	3
s2	1
s3	2
s4	0

student s4 takes no courses

The SELECT COUNT-expression method

SELECT S.sid, (SELECT COUNT(E.Cno) AS NumberCourses   ← Map phase
          FROM   Enroll E
          WHERE E.Sid = S.Sid)   ← Partition phase
FROM   Student S

- Observe that the subquery identified by S.Sid appears in the outer SELECT clause.
- The COUNT of the result of this subquery is then delivered as a value in the outer SELECT clause.
- Notice that this expression must appear between parentheses.
- The output of this query is the same as that on the previous slide.

Example query

- "Find the sid of each student who takes the most courses."
- Using the COUNT FUNCTION method:

```
SELECT  S.sid
FROM    Student S
WHERE NumberOfCourses(S.sid) >= ALL (SELECT NumberOfCourses(S1.Sid)
                                     FROM   Student S1);
```

Example query

- "Find the sid of each student who take the most courses."
- Using the GROUP BY method:

```
WITH
  NumberOfCoursesbyStudent AS (SELECT E.Sid, COUNT(E.Cno) As NumberOfCourses
                               FROM   Enroll E
                               GROUP BY(E.Sid))

SELECT P.sid
FROM   NumberOfCoursesbyStudent P
WHERE  P.NumberOfCourses >= ALL (SELECT P1.NumberOfCourses
                                 FROM   NumberOfCoursesbyStudent P1);
```

Example query

- "Find the sid of each student who takes the most courses."

- Using the COUNT expression method:

```
SELECT   S.Sid
FROM     Student S
WHERE   (SELECT COUNT(E.cno)
          FROM     Enroll E
          WHERE   E.sid = S.sid) >= ALL (SELECT (SELECT COUNT(E.cno)
                                                  FROM   Enroll E
                                                  WHERE   E.sid = S1.sid)
                                          FROM   Student S1);
```

The COUNT-bug of GROUP BY

- The result of the following 2 queries is the same. Notice that there is a bug since, if a student sid takes no courses, then (sid, 0) does not appear in the output.

SELECT E.Sid, COUNT(E.Cno)
FROM   Enroll E
GROUP BY (E.Sid)

SELECT S.Sid, Count(E.Cno)
FROM   Student S, Enroll E
WHERE S.Sid = E.Sid
GROUP BY(S.Sid)

- These two queries give the same result and exhibit the COUNT bug: the tuple (s4,0) does not appear in the result.
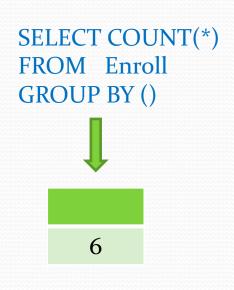
Fixing the COUNT-bug

- To fix the COUNT-bug we need to add the (s,0) pair if student with sid s takes no courses. This can be done with the UNION operator.
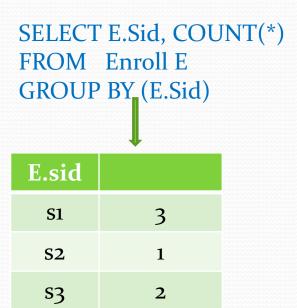
```
(SELECT E.Sid, COUNT(E.Cno) AS No_Courses
 FROM   Enroll E
 GROUP BY (E.Sid))
UNION
(SELECT S.Sid, 0 AS No_Courses
 FROM   Student S
 WHERE S.Sid NOT IN (SELECT E.Sid
                       FROM   Enroll E))
```

Partitioning on different dimensions

Enroll

Sid	Cno
S1	C1
S1	C2
S1	C3
S2	C3
S3	C3
S3	C4

SELECT COUNT(*)
FROM   Enroll
GROUP BY ()

6

SELECT E.Sid, COUNT(*)
FROM   Enroll E
GROUP BY (E.Sid)

E.sid	
S1	3
S2	1
S3	2

SELECT E.Sid, E.Cno, COUNT(*)
FROM   Enroll E
GROUP BY (E.Sid, E.Cno)

E.Sid	E.Cno	
S1	C1	1
S1	C2	1
S1	C3	1
S2	C3	1
S3	C3	1
S3	C4	1

What can appear in the GROUP BY clause?

- Answer: any valid expression over the tuples in the FROM clause.

  SELECT ...
  FROM R1 t1,...,Rn tn
  WHERE ...
  GROUP BY(expression(t1,...,tn))

- Partition: there will be as many cells in the partition as there are different values for expression(t1,...,tn)

Example: expressions in GROUP BY

S

X	Y
2	3
1	3
2	1
0	3

SELECT s.x + s.y AS sum, COUNT(*) as cell_size
FROM   S s
GROUP BY (s.x + s.y)

sum	cell_size
3	2
4	1
5	1

Notice that there are 4 tuples in R assigned to variable s but only 3 different s.x+s.y values: 3, 4, 5. Thus there are only 3 cells in the partition.

Example: expression in GROUP BY

Person

Pid	Age
p1	10
p2	9
p3	12
p3	9

SELECT   p.age > 10 AS OlderThanTen, COUNT(*)
FROM   Person p
GROUP BY (p.age > 10);

OlderThanTen	Count
f	3
t	1

Restrictions on SELECT clause in GROUP BY query

- In a GROUP BY query, the SELECT clause may only contain aggregate expressions that returns a single value for each cell of the partition induced by the GROUP BY clause.

```
SELECT  AggregateExpression(cell(expression(t1,...,tn)), ...
FROM    R t1,...,tn
WHERE condition(t1,...,tn)
GROUP BY (expression(t1,...,tn))
```

Aggregate expressions in SELECT clause

S

X	Y
2	3
1	3
2	1
0	3

SELECT s.x+s.y AS sum,
      SUM(s.x*s.y) AS
          sum_of_products

FROM   S s
GROUP BY (s.x+s.y)

sum	sum_of_products
3	$2 = (2 * 1 + 0 * 3)$
4	$3 = (1 * 3)$
5	$6 = (2 * 3)$

Aggregate expressions in SELECT clause

- The following query will raise an error since s.x is not necessarily unique in a cell defined by s.x+s.y values

S

X	Y
2	3
1	3
2	1
0	3

SELECT s.x
FROM   S s
GROUP BY (s.x+s.y)

The HAVING clause in GROUP BY queries

- The HAVING clause in a GROUP BY selects those cells from the partition induced by the GROUP BY clause that satisfy an Aggregate Condition.
- Only those cells are passed onto the SELECT clause.

```
SELECT AggregateExpression(Cell(expression(t1,…,tn)), …
FROM    R1 t1, …, R tn
WHERE   condition(t1,…,tn)
GROUP BY  (expression(t1,…,tn))
HAVING AggregateCondition(Cell(expression(t1,…,tn)))
```

Example: HAVING clause

- "For each student who majors in CS determine the number of courses taken by that student, provided that this number is at least 2."

SELECT E.Sid, COUNT(E.Cno)
FROM   Enroll E, Student S
WHERE E.Sid = S.Sid AND S.Major = 'CS'
GROUP BY (E.Sid)
HAVING COUNT(E.Cno) ≥ 2;

Simulating HAVING clause with user-defined functions in WHERE

- "For each student who majors in CS, determine the number of courses taken by that student, provided that this number is at least 3."

- The HAVING condition can be simulated in the WHERE clause with user-defined functions.

```
SELECT   S.Sid AS Sid, NumberOfCourses(S.Sid)
FROM     Student S
WHERE    S.major = 'CS' AND NumberOfCourses(S.Sid) ≥ 3
```

Spreadsheet (Data Cube)

Sid	Cno
S1	C1
S1	C2
S1	C3
S2	C1
S2	C2
S3	C2
S4	C1

	c1	c2	c3		sum(sid)	
s1	1	1	1		3	
s2	1	1	0		2	
s3	0	1	0		1	
s4	1	0	0		1	
sum(cno)	3	3	1			
					7	sum()

GROUPING sets

- It may be desirable to simultaneously generate different partitions and then apply an aggregation
- This is supported in SQL via GROUPING sets.

Sid	Cno
S1	C1
S1	C2
S2	C1
S2	C2
S3	C1

→

Sid	Cno	Count
S1	*	2
S2	*	2
S3	*	1
*	C1	3
*	C2	2

SELECT Sid, Cno, COUNT(*)
FROM   Enroll
GROUP BY
  GROUPING SETS((Sid),(Cno))