

Database Programming in an Object-Relational SQL Procedural Programming Language

**PL/pgSQL - Procedural
Language/PostgreSQL**

Dirk Van Gucht¹

¹Indiana University

March 30, 2019

Motivation

- Thus far, we have covered SQL and Object-Relational SQL (abbreviated as OR-SQL) as a language in which each statement correspond to a **single** query, a **single** update, a **single** declaration (definition), etc.
 - That means that a client application must send each query to the database server, wait for it to be processed, receive and process the results, do some computation, then send further queries to the server.
 - All this incurs interprocess communication and will also incur network overhead if the client is on a different machine than the database server
-

Motivation

- In this lecture, we want to address writing **applications**, i.e., **programs**, wherein multiple OR-SQL can be bundled and processed using control statements such as `assignment statements`, `conditional statements`, `loop statements`, etc
 - This correspond to writing programs in an imperative programming languages with the ability to use OR-SQL statements as **embedded code**
 - We will use the PostgreSQL's `plpgsql` language to write such programs
-

Overview

- We begin with a general overview of the programming environment of `plpgsql`
- We give a formal definition of a `plpgsql` program and show how to run it in the PostgreSQL interpreter
- We will illustrate most of this using examples

plpgSQL (Declaration Statements)

Type declarations

CREATE TYPE

Relation declarations

CREATE TABLE

View declarations

CREATE VIEW

Functions declarations

CREATE FUNCTION

Triggers declarations

CREATE TRIGGER

Program variable declarations

DECLARE

Iterator declaration

FOR LOOP, FOREACH

Cursor declaration

DECLARE CURSOR

plpgSQL Garbage collection statements

Type	DROP TYPE
Relation	DROP TABLE
View	DROP VIEW
Function	DROP FUNCTION
Trigger	DROP TRIGGER
Cursor declaration	CLOSE CURSOR
Program variable declarations	not required
Iterator declarations	not required

Expressions and Statements

Expression	Any valid OR-SQL expression including SELECT FROM WHERE expression
Assignment statement	<code>variable_name := expression</code>
Assignment statement	<code>SELECT INTO variable_name</code>
Update statement	<code>INSERT, DELETE, UPDATE</code>
Return statement	<code>RETURN expression</code>
Return query statement	<code>RETURN QUERY query</code>
Function call	<code>SELECT function(parameters)</code>
Block statement	<code>BEGIN ... END</code>
Loop statement	<code>LOOP, WHILE, FOR</code>
Conditional statements	<code>IF ELSE, CASE</code>
Cursor operations	<code>OPEN, FETCH</code>

plpgsql (Program)

- The syntax of a `plpgsql` program is as follows:

```
CREATE OR REPLACE FUNCTION functionName (list of arguments)
    RETURNS return type AS
$$
<label>
<DECLARE declarations>
BEGIN
    sequence of statements;
END <label>;
$$ LANGUAGE plpgsql;
```

- A program is executed in the PostgreSQL interpreter using the call

```
SELECT functionName(parameters);
```


Program (Example with conditional if-else statement)

An example program with the **IF** statement

```
CREATE OR REPLACE FUNCTION convert(a char)
  RETURNS float AS
$$
BEGIN
  IF (a = 't') THEN RETURN 1;
  ELSE
    IF (a = 'f') THEN RETURN 0;
    ELSE
      IF (a = 'u') THEN RETURN 0.5;
      ELSE RETURN(2);
    END IF;
  END IF;
END IF;
END;
$$ LANGUAGE plpgsql;
```

→

SELECT convert('u');
<u>convert</u>
1/2
SELECT convert('z');
<u>convert</u>
2

Program (Example with conditional case statement)

An example program with the **CASE** statement

```
CREATE OR REPLACE FUNCTION convert(a char)
  RETURNS float AS
$$
BEGIN
  CASE WHEN (a = 't') THEN RETURN 1;
        WHEN (a = 'f') THEN RETURN 0;
        WHEN (a = 'u') THEN RETURN 0.5;
        ELSE RETURN 2;
  END CASE;
END;
$$ LANGUAGE plpgsql;
```

→

SELECT convert('u');
<u>convert</u>
0.5
SELECT convert('z');
<u>convert</u>
2

Program (with loop statement)

Iterative program for the **factorial(n)** function

```
CREATE OR REPLACE FUNCTION factorial_iterative (n integer)
RETURNS integer AS
$$
DECLARE
    result integer;
    i integer;
BEGIN
    result := 1;
    FOR i IN 1..n
        LOOP
            result := i * result;
        END LOOP;
    RETURN result;
END;
$$ language plpgsql;
```

Program (with recursion)

Recursive program for the **factorial(n)** function

```
CREATE OR REPLACE FUNCTION factorial_Recursive (n integer)
RETURNS integer AS
$$
BEGIN
    IF n = 0 THEN
        RETURN 1;
    ELSE
        RETURN n * factorial_Recursive(n-1);
    END IF;
END;
$$ language plpgsql;
```

Functions that affect the database state

- Functions can be defined to affect (change) the database state
- Often such functions do not need to return values: they have the **VOID** return type

```
CREATE OR REPLACE FUNCTION change_db_state()  
  RETURNS VOID AS  
$$  
BEGIN  
  DROP TABLE foo_relation;  
  CREATE TABLE foo_relation(a integer);  
  INSERT INTO foo_relation VALUES (1), (2), (3);  
  DELETE FROM foo_relation WHERE a=1;  
END;  
$$ language plpgsql;
```

```
select change_db_state();  
change_db_state  
-----  
→  select * from foo_relation;  
      a  
-----  
      2  
      3
```

Program with local functions

- You can also **CREATE local functions**
- Care must be taken with function **delimiters**

```
CREATE OR REPLACE FUNCTION globalFunction()  
  RETURNS void AS  
$proc$  
BEGIN  
  CREATE OR REPLACE FUNCTION localFunction()  
    RETURNS integer AS  
    $$  
      SELECT 5;  
    $$ language sql;  
END;  
$proc$ language plpgsql;
```

→

```
SELECT globalFunction();  
-----  
globalfunction  
  
SELECT localFunction();  
-----  
localfunction  
5
```

- Notice that **localFunction()** persists after the **SELECT globalFunction()** call

Two kinds of assignment statements

- The typical assignment statement is of the form
 $x := \text{expression};$
- An assignment to a **variable** can also be done with a query and the clause
 $\text{SELECT tuple component(s) INTO variable (s) FROM ... WHERE};$
- The value of the tuple component (s) is (are) assigned to the **variable(s)**

```
CREATE OR REPLACE FUNCTION size_of_A()  
  RETURNS integer AS  
$$  
  DECLARE counter integer;  
  BEGIN  
    SELECT INTO counter COUNT(*) from A;  
    RETURN counter;  
  END;  
$$ language plpgsql
```

→

SELECT * FROM A;
<hr/>
X
'A'
'B'
<hr/>
SELECT size_of_A();
<hr/>
size_of A
<hr/>
2

Special alternative for SELECT INTO assignment statement

```
CREATE OR REPLACE FUNCTION size_of_A()  
  RETURNS integer AS  
$$  
  DECLARE counter integer;  
  BEGIN  
    SELECT INTO counter COUNT(*) from A;  
    RETURN counter;  
  END;  
$$ language plpgsql
```

Since the expression (SELECT COUNT(*) FROM A) evaluates to a single integer, this program can also be written as

```
CREATE OR REPLACE FUNCTION size_of_A()  
  RETURNS integer AS  
$$  
  DECLARE counter integer;  
  BEGIN  
    counter := (SELECT COUNT(*) from A);  
    RETURN counter;  
  END;  
$$ language plpgsql
```


SELECT INTO (non-deterministic behavior)

- SELECT INTO can lead to **non-deterministic (random)** effects!
- This is because SELECT INTO **chooses the first available tuple from the result of the query** and assigns it to the INTO variable (in our case the variable `element_from_A`).³
- Of course, this can be useful when **sampling** data

```
CREATE OR REPLACE FUNCTION choose_one_from_A()
  RETURNS text AS
$$
DECLARE element_from_A text;
BEGIN
  SELECT INTO element_from_A a.x
  FROM (SELECT x from A ORDER BY random()) a;
  RETURN element_from_A;
END;
$$ language plpgsql
```

→

```
SELECT choose_one_from_A();
choose_one_from_a
'B'

SELECT choose_one_from_A();
choose_one_from_a
'A'

SELECT choose_one_from_A();
choose_one_from_a
'A'
```

³If the query does not return any tuple, then the variable is set to NULL.

"Assignment" statements to relation variables

- "Assignment" statements to relation (table) variables are done using the **INSERT INTO**, **DELETE FROM**, and **UPDATE** statements, or using triggers

```
CREATE OR REPLACE FUNCTION relation_assignment()
  RETURNS void AS
$$
BEGIN
  CREATE TABLE IF NOT EXISTS AB(A integer, B integer);
  DELETE FROM AB;
  INSERT INTO AB VALUES (0,0);
  INSERT INTO AB SELECT a1.x, a2.x FROM A a1, A a2;
  UPDATE AB SET A = A*A WHERE B = 2;
END;
$$ language plpgsql;
```

→

```
select * from A;
      x
-----
      1
      2

SELECT * FROM AB;
ERROR: relation "ab" does not exist

SELECT relation_assignment();
SELECT * from AB;
   a  b
----
0    0
1    1
2    1
1    2
4    2
```

Iterators over collections

- Relations and arrays are **collections**
- Relations are unordered collections whereas arrays are ordered collections
- We consider **iterator variables** that slide (move; iterate) over such a collection **one element at a time**
- In SQL, an iterator variable over a relation (which may or may not be the result of a query) is often referred to as a **CURSOR**
- In SQL, it is frequently not necessary to use cursors as the following function illustrates

```
CREATE OR REPLACE FUNCTION there_is_book_that_cost_more_than(k integer)
  RETURNS boolean AS
$$
  SELECT EXISTS(SELECT * FROM book WHERE price > k);
$$ language sql
```

Iterators over collections (**cursors**)

```
CREATE OR REPLACE FUNCTION there_is_book_that_cost_more_than(k integer)
  RETURNS boolean AS
$$
BEGIN
  SELECT EXISTS(SELECT * FROM book WHERE price > k);
END;
$$ language sql;
```

The following function with the same semantics does use the iterator record variable (cursor) **b**

```
CREATE OR REPLACE FUNCTION there_is_book_that_cost_more_than(k integer)
  RETURNS boolean AS
$$
DECLARE exists_book boolean;
      b RECORD; -- the structure will be defined during the program
BEGIN
  exists_book := false;
  FOR b IN SELECT * FROM book -- RECORD b will have have the attribute structure of the book relation
  LOOP
    IF b.price > k
    THEN exists_book := true;
    EXIT;
    END IF;
  END LOOP;
  RETURN exists_book;
END; $$ language plpgsql;
```

Iterators over arrays

- Below is an example from the PostgreSQL manual illustrating iteration through an array using the **FOREACH** clause
- The function **sum** takes an integer array as input and returns the sum of its elements
- The variable **x** is the iterator which gets assigned, one at a time, to each **element** in the array
- Note in particular that **x** is **not assigned to index positions** of the array

```
CREATE FUNCTION sum(A int[])
  RETURNS int8 AS
  $$ DECLARE
    s int8 := 0;
    x int;
  BEGIN
    FOREACH x IN ARRAY A
    LOOP
      s := s + x;
    END LOOP; RETURN s;
  END;
  $$ LANGUAGE plpgsql;
```

Iterators over arrays

On the right is an alternative version for the `sum` function. There an index variable `i` is used that iterates over the index positions of the array.

```
CREATE FUNCTION sum(A int[])
  RETURNS int8 AS
$$ DECLARE
  s int8 := 0;
  x int;
BEGIN
  foreach x IN ARRAY A
  LOOP
    s := s + x;
  END LOOP;
RETURN s;
END;
$$ LANGUAGE plpgsql;
```

↔

```
CREATE FUNCTION sum(A int[])
  RETURNS int8 AS
$$ DECLARE
  s int8 := 0;
  i int;
BEGIN
  FOR i IN array_lower(A,1)..array_length(A,1)
  LOOP
    s := s + A[i];
  END LOOP;
RETURN s;
END;
$$ LANGUAGE plpgsql;
```