


Indexing (B+ Trees)

based on lecture notes by Hector
Garcia-Molina

Indexing




Index and Data File

- **Index** and **Data File** are separate data structures
- **Index** is much smaller than **Data File**. $|\text{Index}| \ll |\text{Data File}|$
- **Index** is stored in a collection of blocks in secondary memory
- **Data File** stored in a collection of blocks in secondary memory
- A key value **K** in the **Index** references a single or multiple records in the **Data File** with that key value **K**
- Records with the same **K** value are chained in the Data File (possibly in sequential blocks)
- Operations (search, insert, delete) on indexed data file require blocks from **Index** and **Data File** to be moved between primary and secondary memory
- Time complexity measured in terms of number of I/O operations

B+Tree Example


$n=3$

n order of B+ tree




Sample non-leaf

Each node in a B+tree
is a block
So n can be quite large





Sample leaf node:





In textbook's notation

$n=3$

Leaf:



Non-leaf:



Size of nodes: $\left\{ \begin{array}{l} n+1 \text{ pointers} \\ n \text{ keys} \end{array} \right.$ (fixed)

Observe that a pointer is a block address

$$(n + 1) * |blockaddress| + n * |key| \leq blocksize$$

$$n \leq \frac{blocksize - |blockaddress|}{|blockaddress| + |key|} \leq \frac{blocksize}{|blockaddress| + |key|}$$

Example: determination of n

- blocksize = 4096 bytes
- |blockaddress| = 8 bytes
- |key| = 9 bytes
- $n \leq \frac{\text{blocksize} - |\text{blockaddress}|}{|\text{blockaddress}| + |\text{key}|}$
- Thus n is maximally 240
- |blockaddress| = 8 bytes permits $2^{32} = 4,294,967,296$ blocks to be referenced
- |blockaddress| = 10 bytes permits 1 trillion blocks to be referenced

Don't want nodes to be too empty

- Use at least

Non-leaf: $\lceil (n+1)/2 \rceil$ pointers

Leaf: $\lfloor (n+1)/2 \rfloor$ pointers to data
+ 1 pointer to next leaf





$n=3$

Non-leaf

Full node

min. node

Leaf



B+tree rules tree of order n

- (1) All leaves at same lowest level
(balanced tree)
- (2) Pointers in leaves point to records
except for “next leaf pointer”
- (3) Root must have at least one key
and two pointers

(3) Number of pointers/keys for B+tree


	Max ptrs	Max keys	Min ptrs	Min keys
Non-leaf (non-root)	$n+1$	n	$\lceil (n+1)/2 \rceil$	$\lceil (n+1)/2 \rceil - 1$
Leaf (non-root)	$n+1$	n	$\lfloor (n+1)/2 \rfloor$	$\lfloor (n+1)/2 \rfloor$
Root	$n+1$	n	2	1

This slide assume that the tree has at least two levels

B+Tree Search

$n=3$

n order of B+ tree



Insert into B+tree

(a) simple case

- space available in leaf


(b) leaf overflow

(c) non-leaf overflow

(d) new root


(a) Insert key = 32

n=3




(a) Insert key = 32

n=3




(a) Insert key = 7

n=3




(a) Insert key = 7

n=3




(a) Insert key = 7

n=3




(c) Insert key = 160

n=3



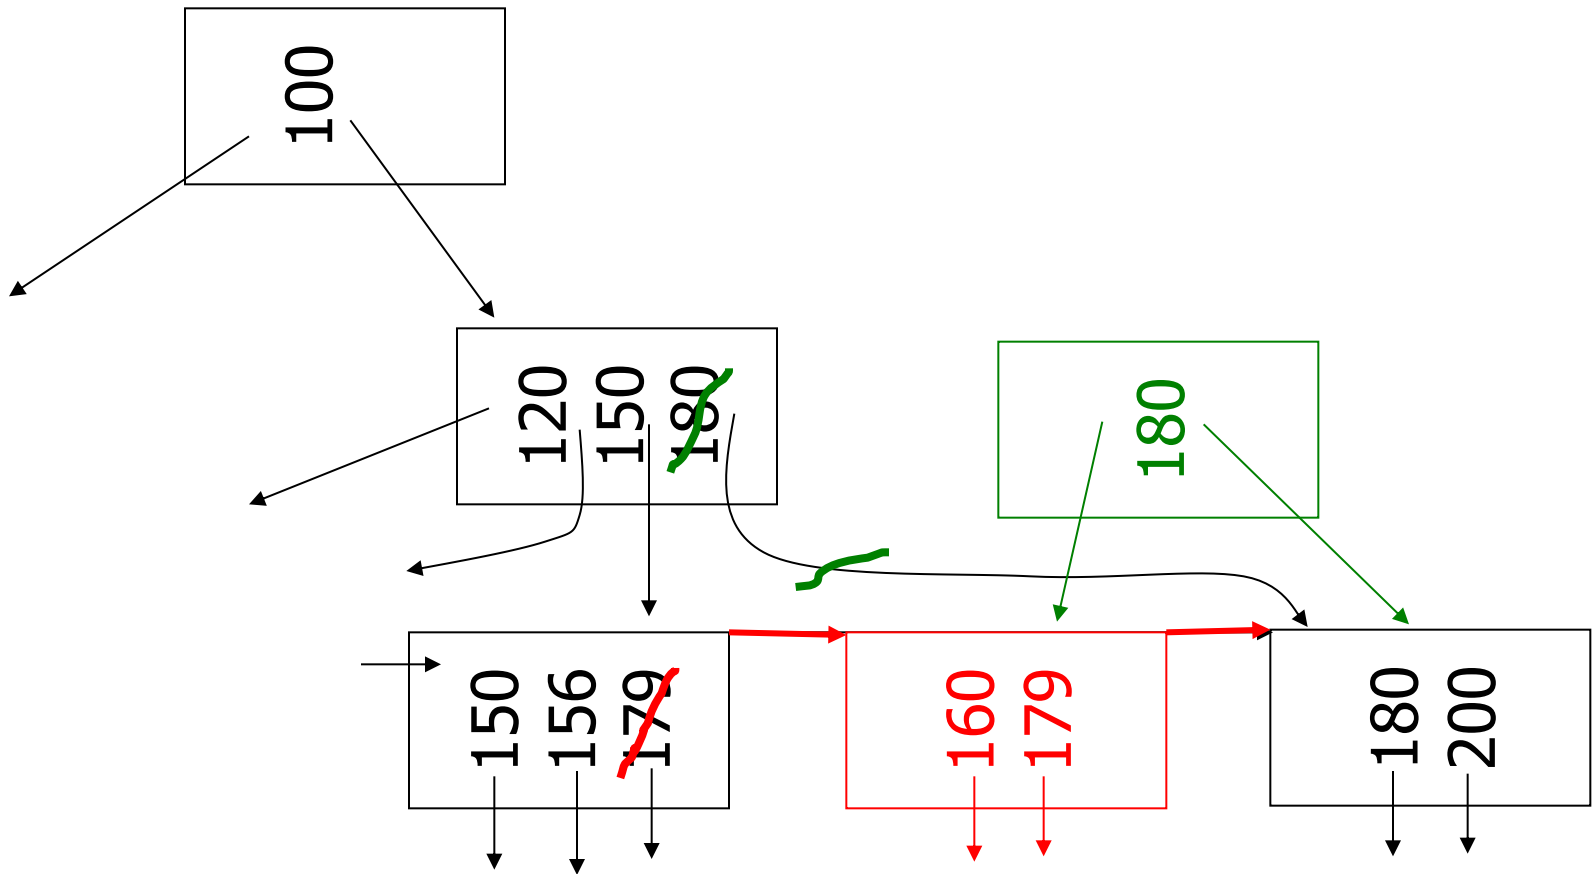
(c) Insert key = 160

n=3



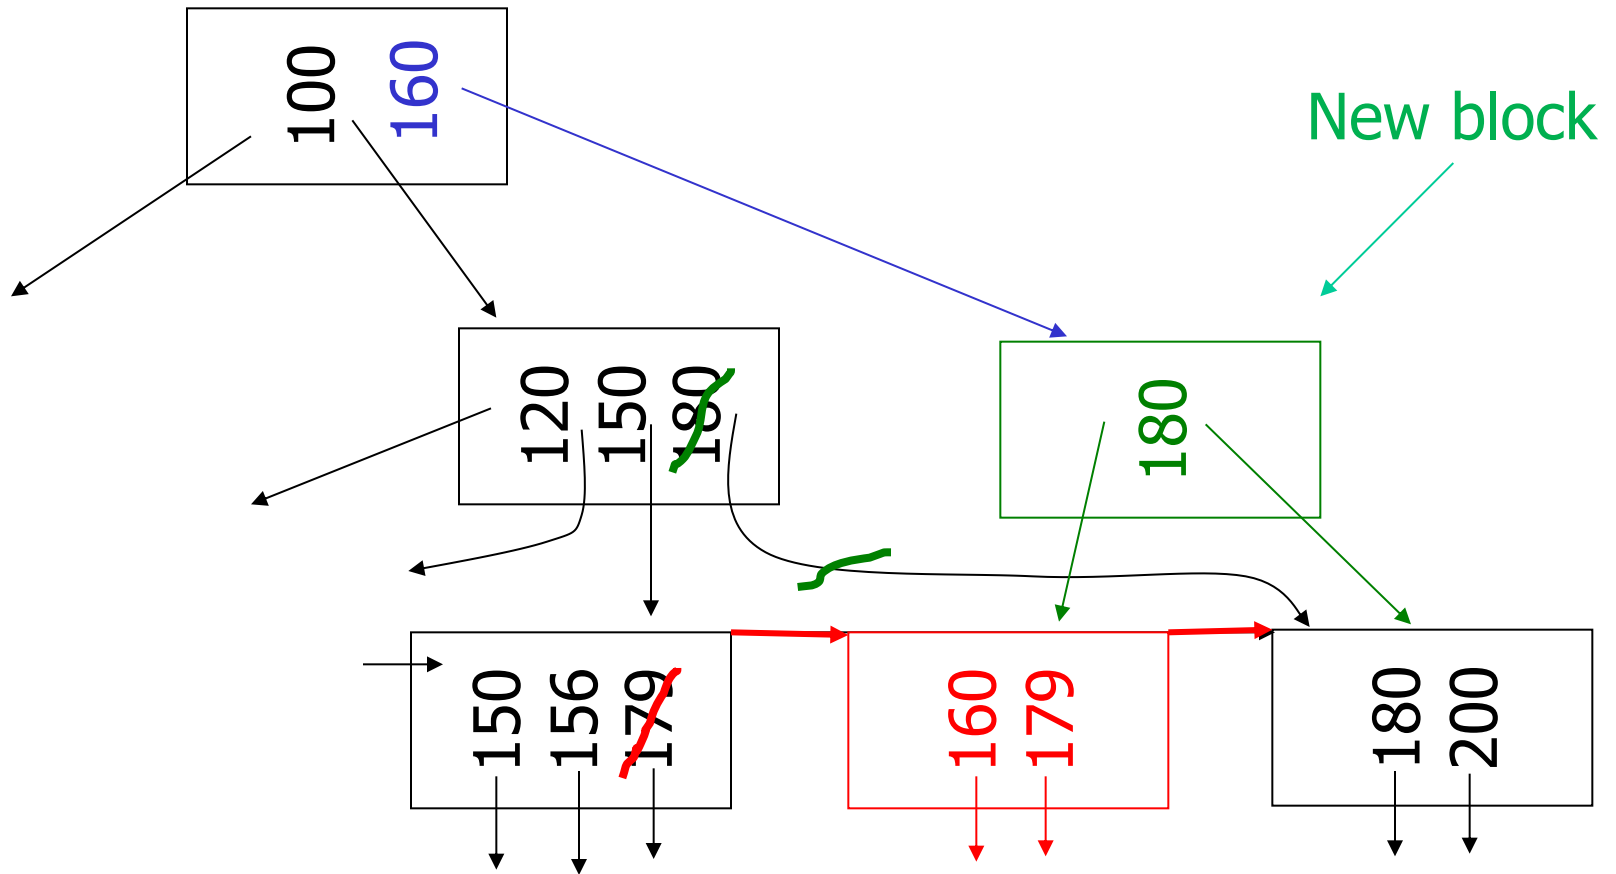
(c) Insert key = 160

n=3




(c) Insert key = 160

n=3




(d) New root, insert 45

n=3




(d) New root, insert 45

n=3




(d) New root, insert 45

n=3



(d) New root, insert 45

n=3




Deletion from B+tree

- (a) Simple case - no example
- (b) Coalesce with neighbor (sibling)
- (c) Re-distribute keys
- (d) Cases (b) or (c) at non-leaf

(b) Coalesce with sibling

– Delete 50


n=4



(b) Coalesce with sibling

– Delete 50


n=4



(c) Redistribute keys

– Delete 50


n=4



(c) Redistribute keys

– Delete 50


n=4



(d) Non-leaf coalesce

– Delete 37


n=4



(d) Non-leaf coalesce

– Delete 37


n=4



(d) Non-leaf coalesce

– Delete 37


n=4



(d) Non-leaf coalesce

– Delete 37

n=4



Complexity analysis of search

Assumption: N blocks in the index
in order or the tree

Each search requires navigating along a path
from the root to a leaf

Thus the complexity corresponds to the
height h of the tree

height h is maximum if the branching factor at the nodes is minimal

at the root: 2

at non-leaf nodes: essentially $n/2$

the root splits the tree into two trees of $N/2$ keys

we must find the number a such that $(\frac{n}{2})^a \geq N/2$

So $a \sim \log_{n/2}(N/2) = O(\log_n(N))$

If the branching factor is maximum at each nodes, the complexity analysis also gives $O(\log_n(N))$

Complexity analysis for insert and delete

In the worst case, for both insert and delete, processing is determined by a downward phase of h steps and an upward phase also of h steps

Consequently, the complexity of insert and delete is also $O(\log_n(N))$

To improve complexity, place the first several levels in main memory

For typical cases, search time is measured in terms of a few block I/O's and this for very large data files

The leaf level provides a sorted list of the records in the data file.

Range searches can be accommodated: given range (k_1, k_2) , locate the leaf holding k_1 and then follow along the leaf level until reaching records with key value higher than k_2