

Algorithms for Relational Algebra Operations (Join Algorithms)

Overview of Join Algorithms

- Nested-loop and block nested-loop join $R \bowtie_{\theta} S$
- Sort-merge join $R \bowtie S$
- Hash-based join strategies $R \bowtie S$
- Index join $R \bowtie S$

Join Operator

- JOIN: Most important relational operator
 - Potentially very expensive
 - Required in all practical queries and applications
 - Often appears in groups of joins
- Example: Relations R (A, B) and S (B, C)

```
SELECT *  
FROM R JOIN S ON (R.B  $\theta$  S.B)  
with  $\theta$ : =,  $\neq$ , <,  $\leq$ , >,  $\geq$ 
```

Unary versus Binary Operations

- Relational operators working on one table
 - Selection, projection
- On two tables
 - Product, Join, Semi-join, Intersection, Union, Difference
- **Binary operators** are usually more expensive
 - Unary: Look at table (scanning, index, hash, ...)
 - Binary: Look at each tuple of first table **for each tuple of** second table
 - “Potentially” quadratic complexity

Nested-loop Join

Nested-loop join

```
FOR EACH  $r$  IN  $R$  DO
  FOR EACH  $s$  IN  $S$  DO
    IF ( $r.B \theta s.B$ ) THEN OUTPUT ( $r \bowtie_{\theta} s$ )
```

Some improvement (block-based)

```
FOR EACH block  $x$  IN  $R$  DO
  FOR EACH block  $y$  IN  $S$  DO
    FOR EACH  $r$  in  $x$  DO
      FOR EACH  $s$  in  $y$  DO
        IF ( $r.B \theta s.B$ ) THEN OUTPUT ( $r \bowtie_{\theta} s$ )
```

- Cost estimations
 - $b(R)$, $b(S)$ number of blocks in R and in S , respectively
 - Each block of outer relation is read once
 - Inner relation is read once for each block of outer relation
 - Inner two loops are free (only main memory operations)
 - Altogether: $b(R) + b(R) * b(S)$

Example

- Assume $b(R)=10,000$, $b(S)=2,000$
- R as outer relation
 - $IO = 10,000 + 10,000 * 2,000 = 20,010,000$
- S as outer relation
 - $IO = 2,000 + 2,000 * 10,000 = 20,002,000$
- Use **smaller relation as outer relation**
 - For large relations, choice doesn't really matter
- Can we do better?

Block nested-loop join

- M – Size of main memory in blocks (or number of buffers)
- Rule of thumb: **Use all memory you can get**
 - Use all memory the buffer manager allocates to the process
- Blocked-nested-loop

```
FOR i=1 TO b(R)/M DO
  READ NEXT R-chunk of M blocks of R into Memory buffer
  FOR EACH block y IN S DO
    FOR EACH r in R-chunk DO
      FOR EACH s in y do
        IF ( r.B  $\theta$  s.B ) THEN OUTPUT (r  $\bowtie_{\theta}$  s)
```
- Cost estimation
 - Outer relation is read once
 - Inner relation is read once for **every chunk** of R
 - There are $\sim b(R)/M$ chunks
 - IO = **$b(R) + b(R)*b(S)/M$**

- Example
 - Assume $b(R)=10,000$, $b(S)=2,000$, $M=500$
 - R as outer relation
 - $IO = 10,000 + 10,000 * 2,000 / 500 = 50,000$
 - S as outer relation
 - $IO = 2,000 + 2,000 * 10,000 / 500 = 42,000$
 - Compare to one-block NL: 20,002,000 IO
- Use smaller relation as outer relation
- But sizes of relations do matter:
 - If one relation fits into memory ($\min(b(R), b(S)) < M$)
 - Total cost: $b(R) + b(S)$

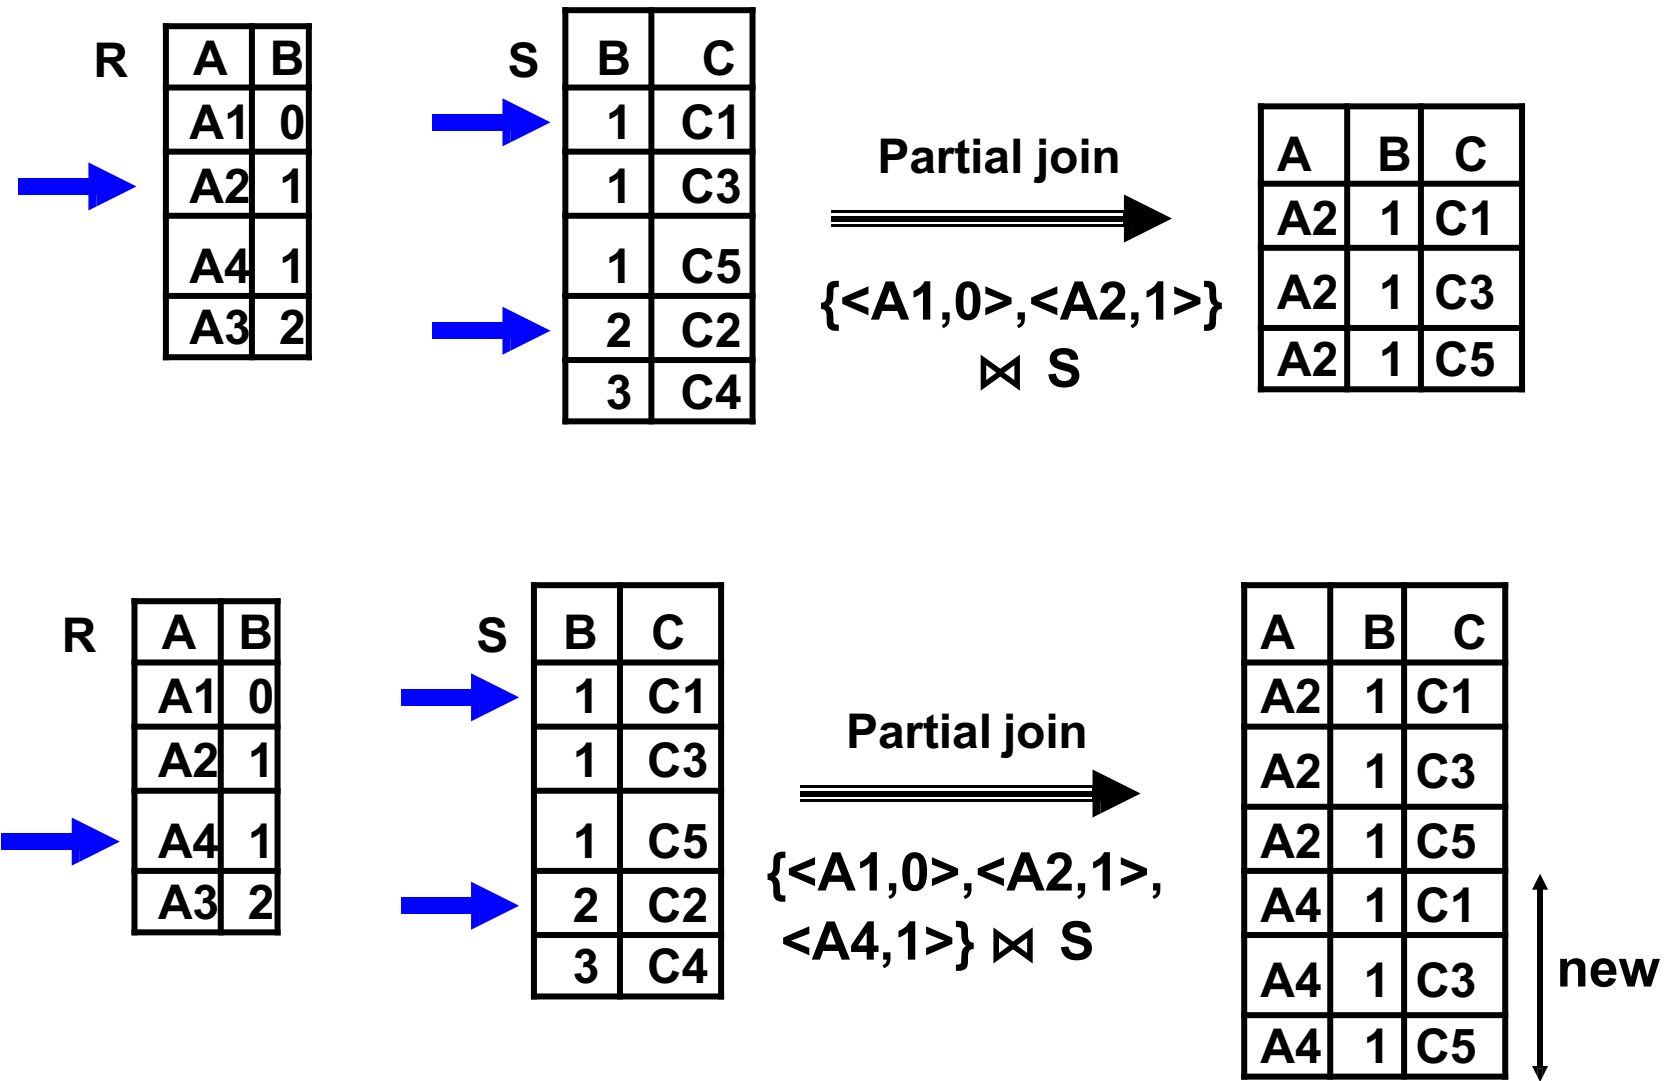
Sort-Merge Join

- How does it work?
- Only works for Natural Join $R \bowtie S$
- What does it cost?
- Does it matter which is outer/inner relation?
- When is it better than block-nested loop?
- Be concerned about skew on join attribute

Sort-Merge Join

- How does it work?
 - Sort both relations on join attribute(s)
 - Merge both sorted relations
- Caution if duplicates (skew) exist
 - The result size still is $|R| * |S|$ in worst case
 - If there are r/s tuples with value x in the join attribute in R / S , we need to output $r*s$ tuples for x
 - More importantly, all these r/s must simultaneously fit in main memory (not always true if there is skew)

Example



Cost estimation (without skew)

- Sorting R costs $2 * b(R) * \text{ceil}(\log_M(b(R)))$
- Sorting S costs $2 * b(S) * \text{ceil}(\log_M(b(S)))$
- Merge phase reads each relation once
- Total IO
 - $b(R) + b(S) + 2 * b(R) * \text{ceil}(\log_M(b(R))) + 2 * b(S) * \text{ceil}(\log_M(b(S)))$

This is only the case when, for each join-value b , the R and S blocks with that value b fit together in the buffer.

If this is not the case for value b , then the sort-merge algorithm needs to do a local block-nested join loop on the R and S blocks with value b .

Better than Block-Nested-Loop?

- Assume $b(R)=10,000$, $b(S)=2,000$, $M=500$
 - BNL costs 42,000
 - With S as outer relation
 - SM: $10,000+2,000+4*10,000+4*2,000 = 60,000$ since
 $\text{ceil}(\log_{500}(10,000)) = \text{ceil}(\log_{500}(2,000)) = 2$
- Assume $b(R)=1,000.000$, $b(S)=1,000$, $M=500$
 - BNL costs $1,000 + 1,000,000*1000/500 = 2,001,000$
 - SM: $1.000.000+1.000+6*1.000.000+4*1.000 = 7,005,000$

Comparison


- Assume relations of equal size: B blocks
- SM: $2*B + 4*B*\log_M(B)$
- BNL: $B+B^2/M$
- BNL > SM
 - $B+B^2/M > 2*B + 4*B*\log_M(B)$
 - $B/M > 1+ 4*\log_M(B)$
 - $B > M + 4*M*\log_M(B)$

Hash Join

- As always, we may save sorting if good hash function available
- Assume a **very good** hash function
 - Distributes hash values **almost uniformly** over hash table
 - If we have **good histograms** (later), a simple interval- based hash function will usually work
- How can we apply hashing to joins?

Hashing a file on join attribute(s)


- Idea: partition a relation R into buckets, on disk
- Each bucket has size approx. $B(R)/M$



- Does each bucket fit in main memory ?
 - Yes if $B(R)/M \leq M$, i.e. $B(R) \leq M^2$

Hash Join Idea

- Use join attributes as hash keys in both R and S
- Choose hash function for **hash table of size M**
 - Each bucket has size $b(R)/M$, $b(S)/M$
- Hash phase
 - Scan R, compute hash table, **writing full blocks to disk immediately**
 - Scan S, compute hash table, **writing full blocks to disk immediately**
- Merge phase
 - Iteratively, load same bucket of R and of S in memory
 - Compute join
- Total cost
 - **Hash phase** costs $2*b(R)+2*b(S)$
 - **Merge phase** costs $b(R) + b(S)$
 - Total: **$3*(b(R)+b(S))$**
 - Under what assumption??



Index Join


- Assume we have an index “**B_Index**” on one join attribute
- Choose indexed relation as inner relation
- Index join

```
FOR EACH r IN R DO
```

```
  X = { SEARCH (S.B_Index, <r.B>) }
```

```
  FOR EACH TID i in X DO
```

```
    s = READ (S, i) ; output (r ⋈ s).
```



S.B_Index auf S

- Actually, this is a one **block-nested loop with index access**
 - Using BNL possible (and better)

Semi Join

- Consider queries such as
 - `SELECT DISTINCT R.* FROM S,R WHERE R.B=S.B`
 - `SELECT R.* FROM R WHERE R.B IN (SELECT S.B FROM S)`
- What's special?
 - No values from S are requested in result
 - S (or inner query) acts as filter on R
- Semi-Join $R \bowtie S$

Implementing Semi-Join

- Using blocked-nested-loop join
 - Choose relation R as outer relation
 - Perform BNL
 - Whenever partner for R.B is found, exit inner loop
- Using sort-merge join
 - Sort R
 - Sort join attribute values from S, remove duplicates on the way
 - Perform merge phase as usual
- Using hash join
 - Hash R
 - Hash join values from S, remove duplicates on the way
 - Perform hash phase as usual

Implementing Intersection, Union, Difference

- Analogous with semi-join
- Using sort-merge join
 - Sort R, remove duplicates on the way
 - Sort S, remove duplicates on the way
 - Perform merge phase as usual (checking for and, or, not)
- Using hash join
 - Hash R, remove duplicates on the way
 - Hash S, remove duplicates on the way
 - Perform hash phase as usual (checking for and, or, not)

Observe that the performance is never $O(R*S)$