

**B461 – Database Concepts
Practice Problems
Fall 2021**

For these practice problems we will use the following relation schema:

- Person (pid, pname, city)
- Company (cname, city)
- Skill (skill)
- jobSkill (pid, skill)
- worksFor (pid, cname, salary)
- Manages (mid, eid)

In this database we maintain a set of persons (Person), a set of companies (Company), and a set of skills (Skill). The city attribute in Person specifies the city in which the person lives. The city attribute in Company indicates a city in which the company is located. (Companies may be located in multiple cities.) A person can be employed by at most one company. (We permit that a person is not employed but that a company must have at least one employee.) A person can have multiple job skills (jobSkills). A job skill can be a skill of multiple persons. (A person may not have any job skills, and a skill may have no person who has that skill.) A pair (m; e) in Manages indicates that person e is a manager of person m. The attributes mid and eid refer to the pids of the manager and the employee, respectively. We permit that a manager manages multiple employees and that an employee can have multiple managers. (It is possible that an employee is not a manager and that an employee has no manager.) We further require that if a manager manages an employee, then that employee must work for the company where that manager is also an employee.)

1. Identify and specify the likely foreign keys in the above relational schema
2. Draw an entity-relationship diagram that models this relational schema. Include the constraints on the relationships. (In the exam, you can either
 - a. Draw the diagram on a paper, take a picture with your phone, transfer the picture to your laptop and upload it on canvas as your response OR
 - b. Write the entities (with attributes), relationships (with attributes and constraints) as your answer

Formulate the following queries in the relational algebra: You can use the abbreviations P, C, S, JS, W, and M for the relations Person, Company, Skill, jobSkill, worksFor, and Manages

3. Find the pid and pname of each person who lives in the same city as one or more of his or her managers.
4. Find the cname of each company that is located in 'Bloomington' but in no other city.

Translate each of the following queries into an equivalent relational algebra expression in RA notation. (You are not required to optimize these RA expressions.) We recommend that you show some intermediate step in the translation. This will permit us to give partial credit for correct translation steps.

5. Query 1

```
SELECT m.mid
FROM Manages m, worksFor w1, worksFor w2
```

```

WHERE m.mid = w1.pid
      AND m.eid = w2.pid
      AND w1.salary < w2.salary
      AND w1.cname = 'Google'

```

6. Query 2


```

SELECT p.pid
FROM Person p, worksFor w
WHERE p.city = 'Indy'
      AND p.city <> w.city
      AND p.id NOT IN (SELECT m.mid
                       FROM manages m
                       WHERE EXISTS (SELECT 1
                                     FROM worksFor w
                                     WHERE m.eid = w.pid
                                     AND w.salary > 20000))

```

Formulate the following queries using the object relational SQL

7. Create a view companyEmployees(cname,employees) which associates with each company, identified by a cname, the set of pids of persons who work for that company.
8. Define a view cityHasCompanies(city,companies) which associates with each city the set of cnames of companies who are located in that city.
9. Define a view personHasSkills(pid,skills) which associates with each person, identified by a pid, his or her set of job skills. Observe that a person may not have any job skills.
10. Consider the following B+-tree of order n=2 that indexes records with keys 3, 8, 10, and 11



Show the contents of your B+-tree index after inserting records with keys 0, 7, 14, and 9 in that order.